



WS-A3 - Network App Development: Design , Deploy , & Automate

Autocon3 - Deck 1
Mau - Deepinder - Emre

A U T O C O N 3

THE NETWORK AUTOMATION CONFERENCE

Who we are?

- Network automation professionals from Nokia
- Diverse backgrounds, generations, and experiences
- United by a passion for learning, building, and improving

Introduction & Foundations

- **Workshop Goals & Key Principles:** Declarative automation, event-driven design (Mau)
- **Technology Overview:** Containerlab, Kind, iPerf, Reconcilers, gNMIc (Mau)
- **Managing Container Images:** Building, posting locally, and pushing to online registries (Emre)
- **Preparing an App for Containers:** Structuring code for easy deployment (Emre)

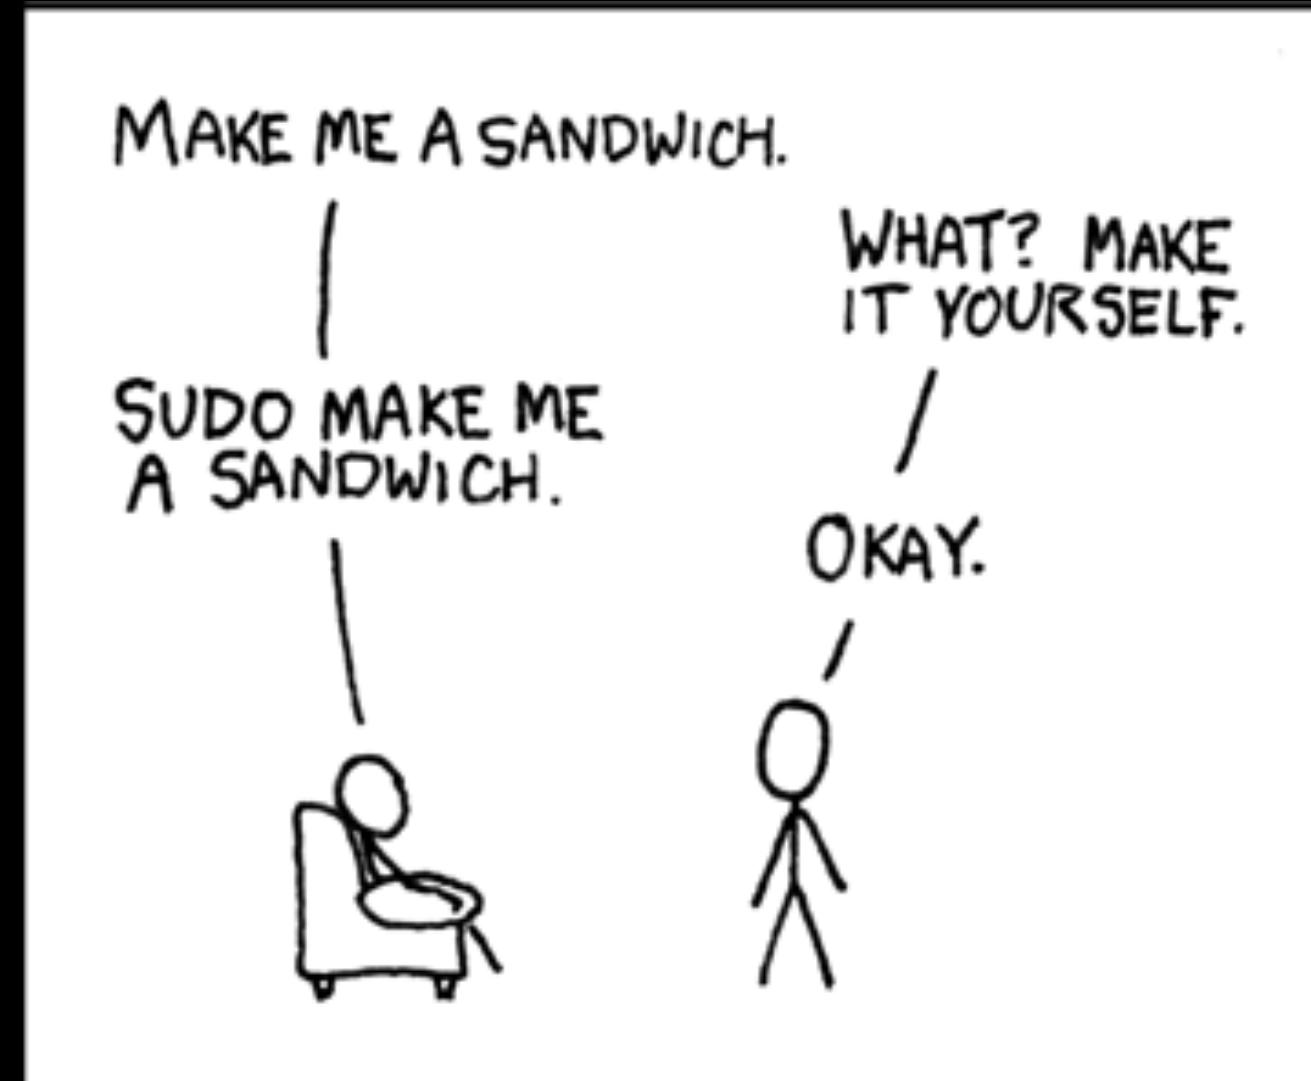
Goals



Goals

- Is there a better way?
- Discover tools that simplify and scale
- Build more with less code
- Learn from use cases and lessons learned
- Apply practical DevOps and automation techniques
- Create resilient, shareable, and maintainable solutions

Declarative vs Imperative

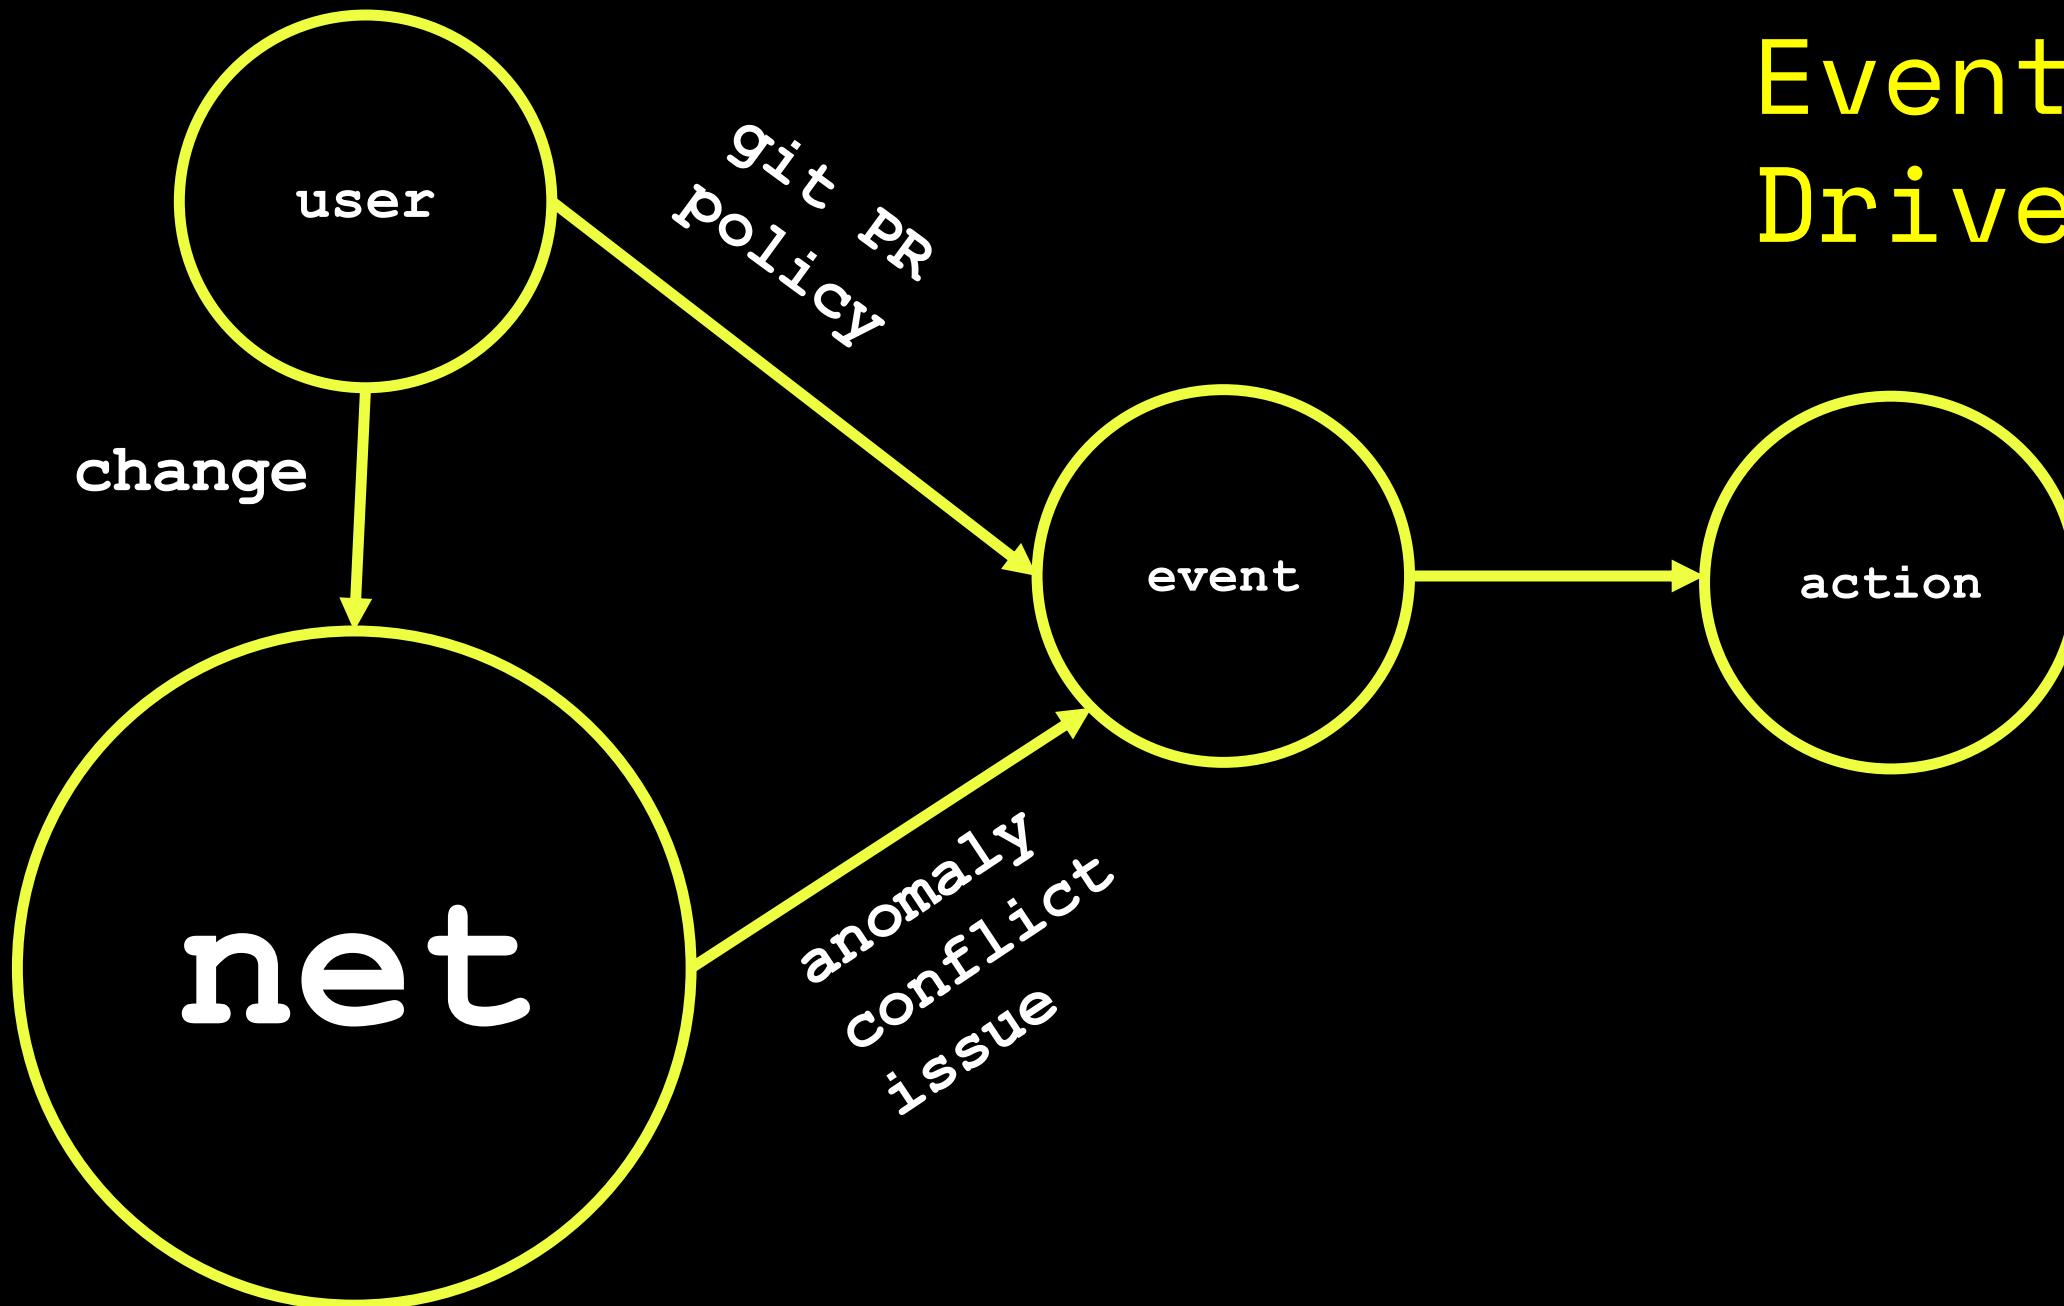


```
- name: Run iperf3 clients to multiple destination ports
hosts: iperf_clients
gather_facts: no
vars:
  iperf_server_ip: "192.168.1.100"
  destination_ports: [5201, 5202, 5203]
  parallel_clients_per_port: 2
Imperative
tasks:
- name: Run iperf3 client instances
  ansible.builtin.shell: |
    nohup iperf3 -c {{ iperf_server_ip }} -p {{ item.1 }} -t 60 >
    /tmp/iperf3_{{ item.0 }}_port{{ item.1 }}.log 2>&1 &
    loop: "{{ query('subelements', [range(1,
parallel_clients_per_port+1), destination_ports]) }}"
  async: 0
  poll: 0
  ignore_errors: yes
```

Declarative

```
apiVersion: example.com/v1
kind: IperfClient
metadata:
  name: iperf3-client
spec:
  targetIP: "172.254.102.101" # Replace with server
  cluster IP
  initPort: 30001
  endPort: 30005
  image: iperf3-client:0.1a
```

Event Driven





Lab Prep

Lab Setup Options

Two Ways to Join the Labs:

- **GitHub Codespaces (Recommended)**
 - No installation needed – runs in your browser
 - Requires a GitHub account and ~4 hours of Codespaces credits
 - Remember to switch to 4 vCPUs
- **Local Deployment**
 - Run everything using **VS Code + Docker Desktop**
 - Requires 16 GB RAM & 8 Cores minimum
 - Watch out for WSL

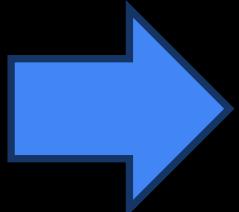
Required Tools & Extensions

Local Setup Checklist (if not using Codespaces):

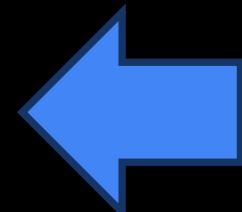
- Install:
 - VS Code
 - Docker Desktop
 - Git
- VS Code Extensions:
 - Dev Containers
 - Remote - SSH

Port Mapping

codespaces



PROBLEMS	OUTPUT	DEBUG CONSOLE	TERMINAL	PORTS 5
Port	Forwarded Address		Runn	
● 2181				https://laughing-eureka-j... /usr/b
● 8001	🔗	×		https://laughing... ↗ ⓘ
● 9092				https://laughing-eureka-j... /usr/b
● 9093				https://laughing-eureka-j... /usr/b
● 9095				https://laughing-eureka-j... /usr/b



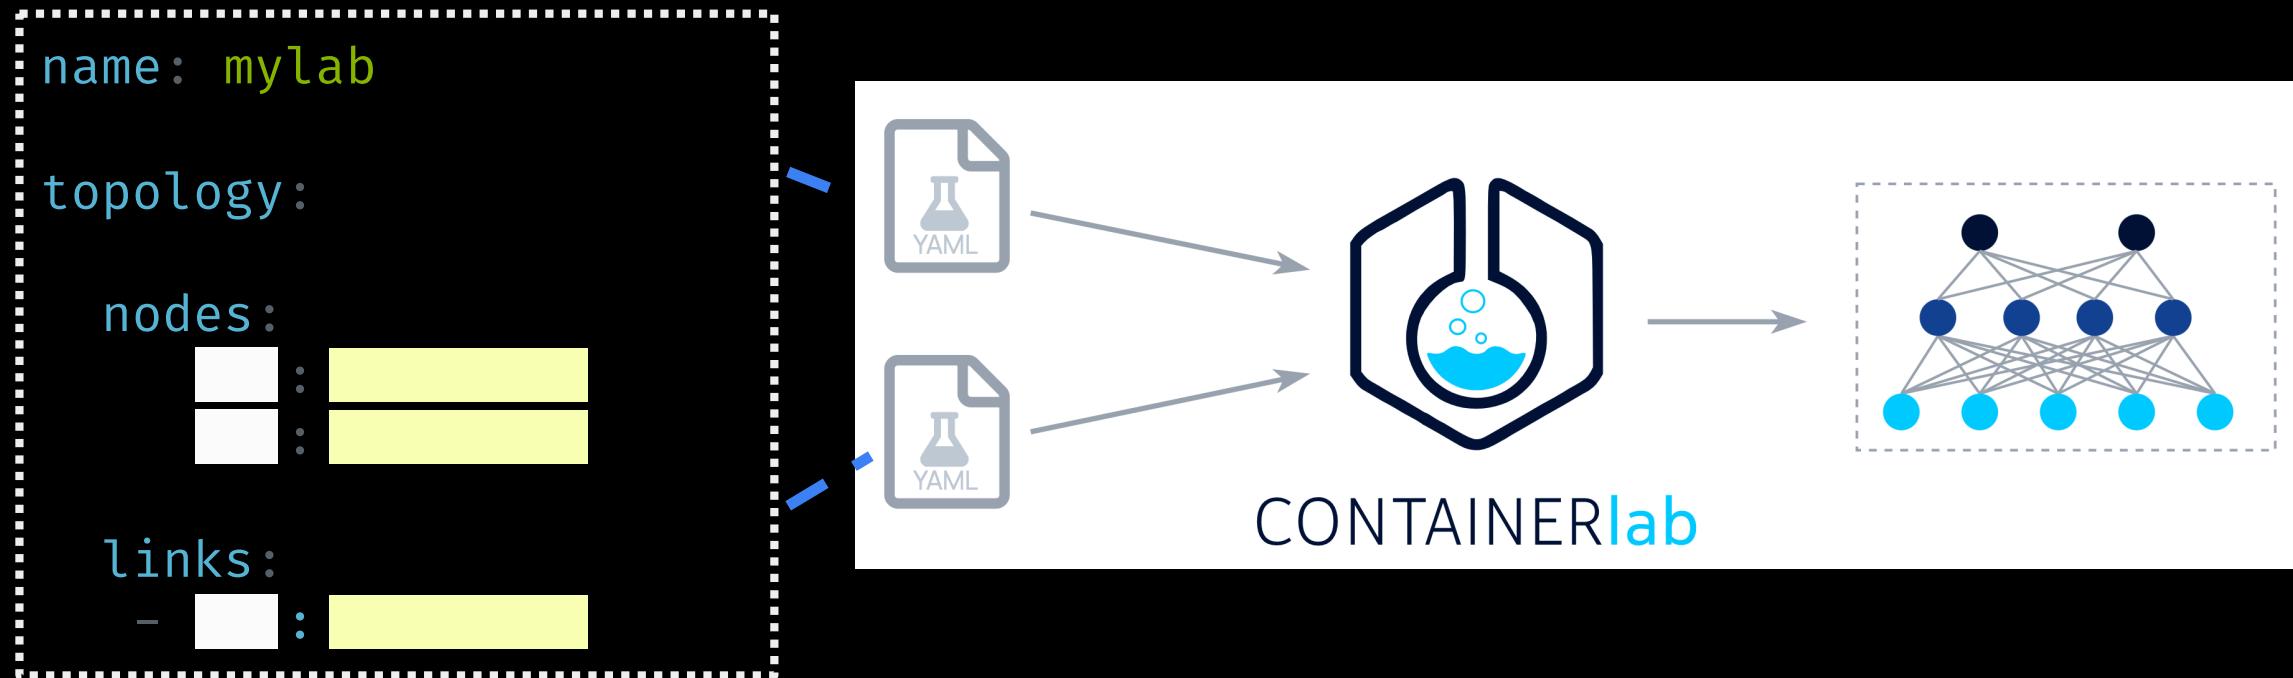
local

PROBLEMS	OUTPUT	DEBUG CONSOLE	TERMINAL	PORTS 13
Port	Forwarded Addr...	Running F		
● 2181	localhost:2181	/usr/bin/do		
● 2222	localhost:2222	sshd: /usr/		
● 3000	localhost:3000	/usr/bin/do		
● 8001	localhost:8001	/usr/bin/do		
● 9092	localhost:9092	/usr/bin/do		
● 9093	localhost:9093	/usr/bin/do		
● 9095	localhost:9095	/usr/bin/do		

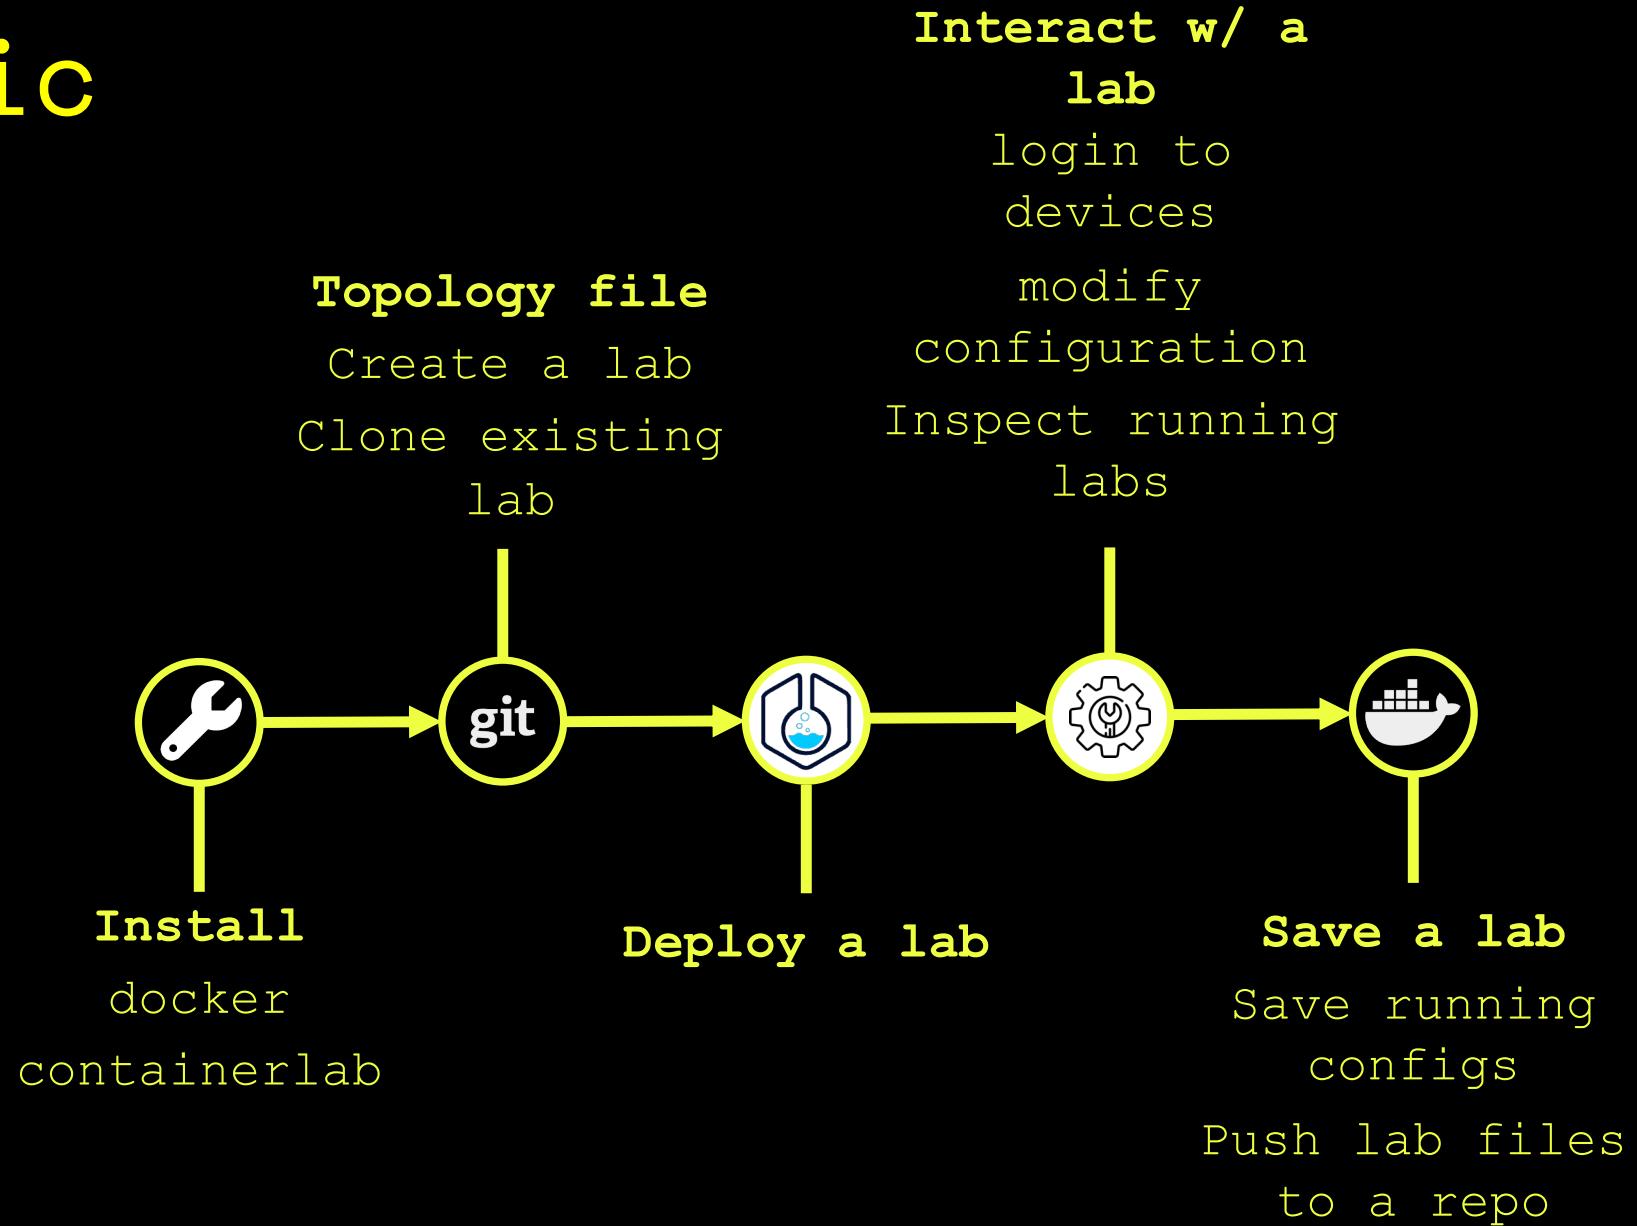


Technology Overview

Containerlab or clab



Clab: basic workflow



Topology File (e.g. topo.yml)

```
topology:  
  kinds:  
    nokia_srlinux:  
      type: ixrd21  
      image: ghcr.io/nokia/srlinux:24.7.2  
  nodes:  
    k8s01:  
      kind: k8s-kind  
    k8s02:  
      kind: k8s-kind  
    k8s01-control-plane:  
      kind: ext-container  
    exec:  
      - "ip addr add 172.254.101.101/24 dev eth1"  
      - "ip route add 172.0.0.0/8 via 172.254.101.1"
```



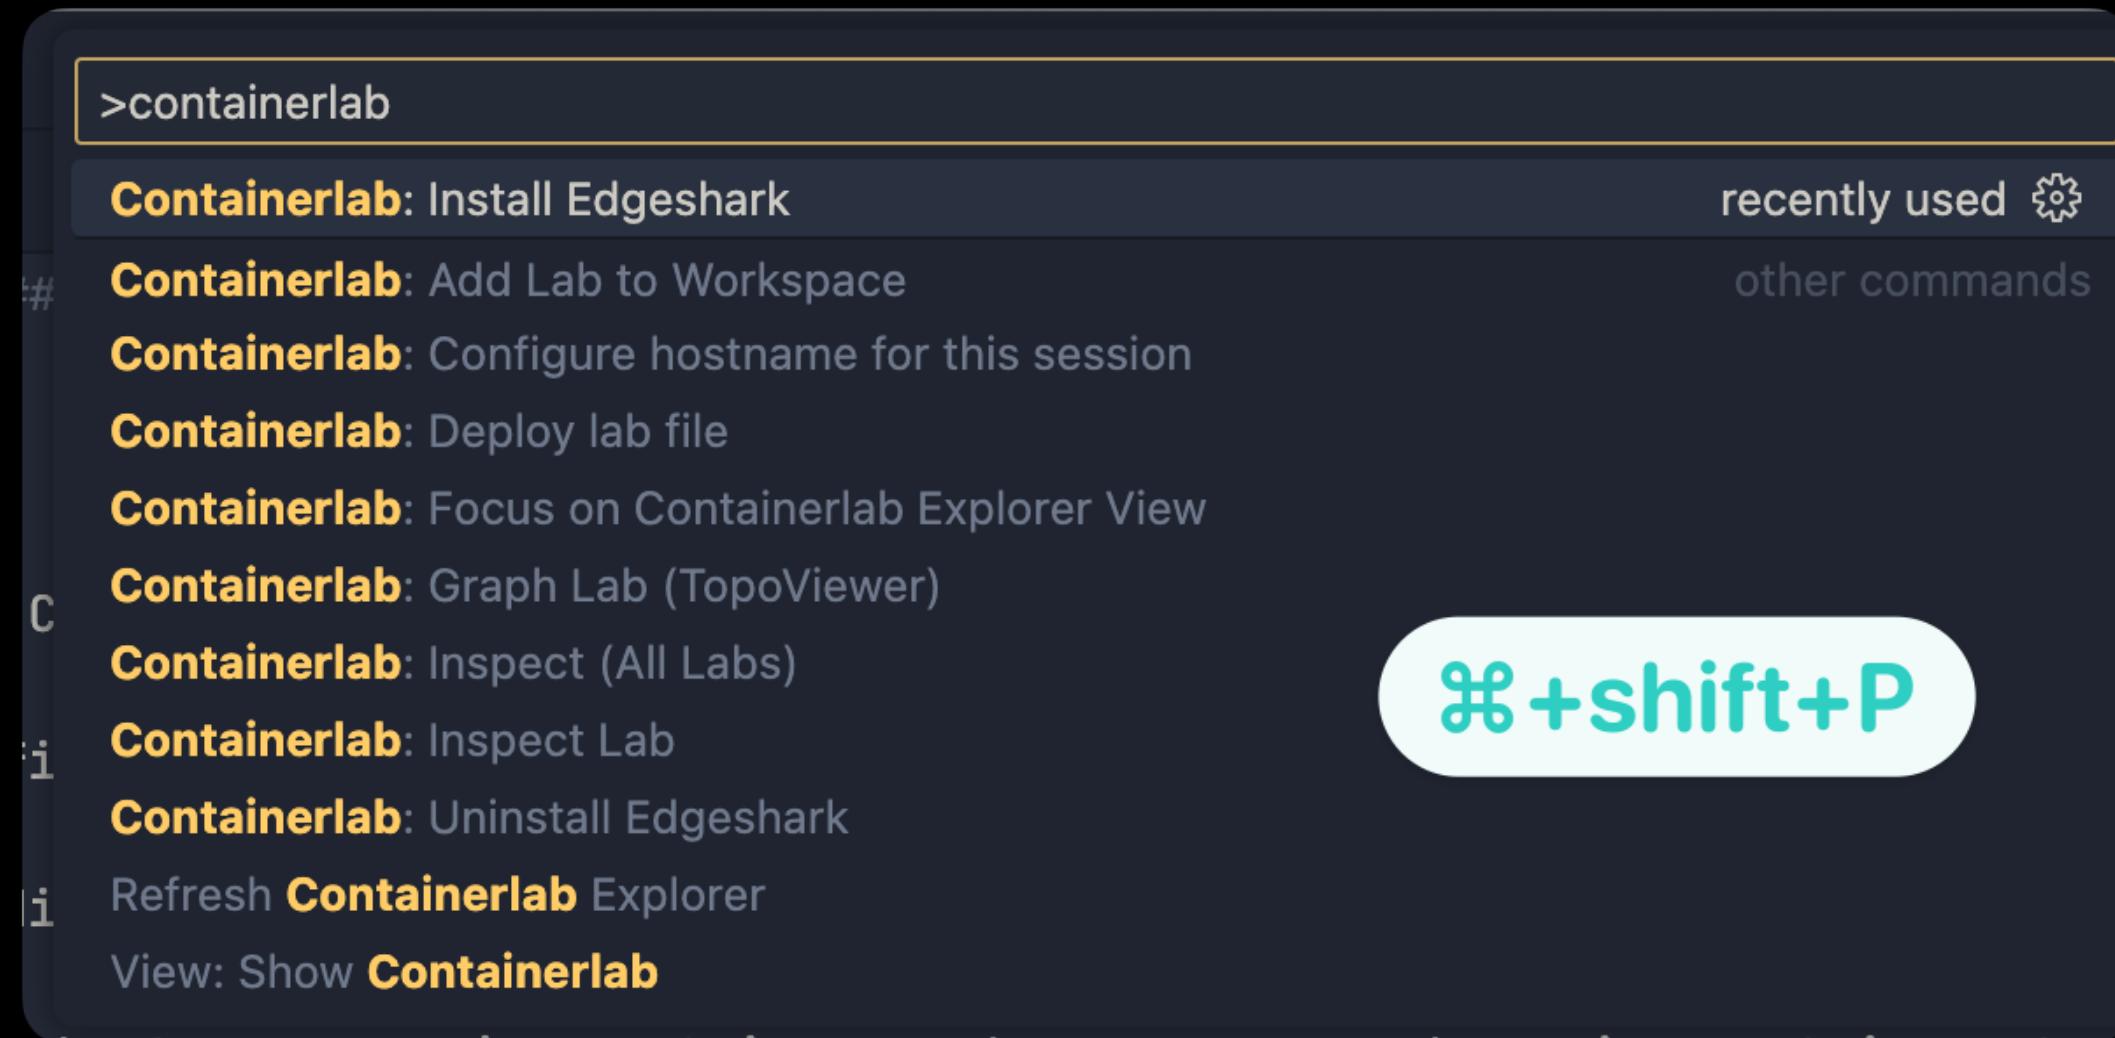
Clab: Deploy a lab

```
[srl-apps]# clab deploy -t topo.yml
INFO[0000] Containerlab v0.25.1 started
INFO[0000] Parsing & checking topology file: topo.yml
INFO[0000] Creating lab directory: /root/pygnmi-srl-apps/clab-dc-k8s
INFO[0000] Creating docker network: Name="kind",
IPv4Subnet="172.18.100.0/16", IPv6Subnet="", MTU="1500"
INFO[0000] Creating container: "grafana"
INFO[0000] Creating container: "SPINE-DC-2"
INFO[0000] Creating container: "prometheus"
INFO[0000] Creating container: "LEAF-DC-1"
INFO[0000] Creating container: "BORDER-DC"
INFO[0000] Creating container: "SPINE-DC-1"
INFO[0000] Creating container: "LEAF-DC-2"
```

Clab: Inspect a lab

```
[srl-apps]# clab inspect -t topo.yml
INFO[0000] Parsing & checking topology file: topo.yml
+-----+-----+-----+-----+
| # |      Name      | Container ID |      Image      | Kine
+-----+-----+-----+-----+
| 1 | clab-dc-k8s-BORDER-DC | 9876f09a5580 | ghcr.io/nokia/srlinux:21.6.4 | sr
| 2 | clab-dc-k8s-LEAF-DC-1 | 830369bb4d39 | ghcr.io/nokia/srlinux:21.6.4 | sr
| 3 | clab-dc-k8s-LEAF-DC-2 | 05d303e50816 | ghcr.io/nokia/srlinux:21.6.4 | sr
| 4 | clab-dc-k8s-SPINE-DC-1 | 574ff19416fb | ghcr.io/nokia/srlinux:21.6.4 | sr
| 5 | clab-dc-k8s-SPINE-DC-2 | e44d29973290 | ghcr.io/nokia/srlinux:21.6.4 | sr
| 6 | clab-dc-k8s-grafana   | e6d5221fa472 | grafana/grafana:latest    | lin
| 7 | clab-dc-k8s-prometheus | 533473420ff1 | prom/prometheus:latest    | lin
+-----+-----+-----+-----+
```

Clab: VSCode



Kubernetes kind

- Local Kubernetes clusters using Docker container “nodes”.
- Designed for testing Kubernetes.
- Requirements:
 - go (1.17+)
 - Uses docker for node instances



Kind: Install

- You can install kind with:
`go install sigs.k8s.io/kind@v0.18.0`
- This will put kind in `$(go env GOPATH)/bin`.
 - You may need to add that directory to your `$PATH` as shown here if you encounter the error kind: command not found after installation.
- Kind uses docker for node instances
- Once you have docker running you can create a cluster with:

`kind create cluster`

Kind: Config

- For this demo we'll use **ONE** controller
- Use “**kind load**” to upload images for apps like **iperf** after cluster is created
 - You can to setup a private registry

```
[~/kind]# cat cluster_datacenter.yaml
---
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
nodes:
  - role: control-plane
  - role: worker
  - role: worker
```



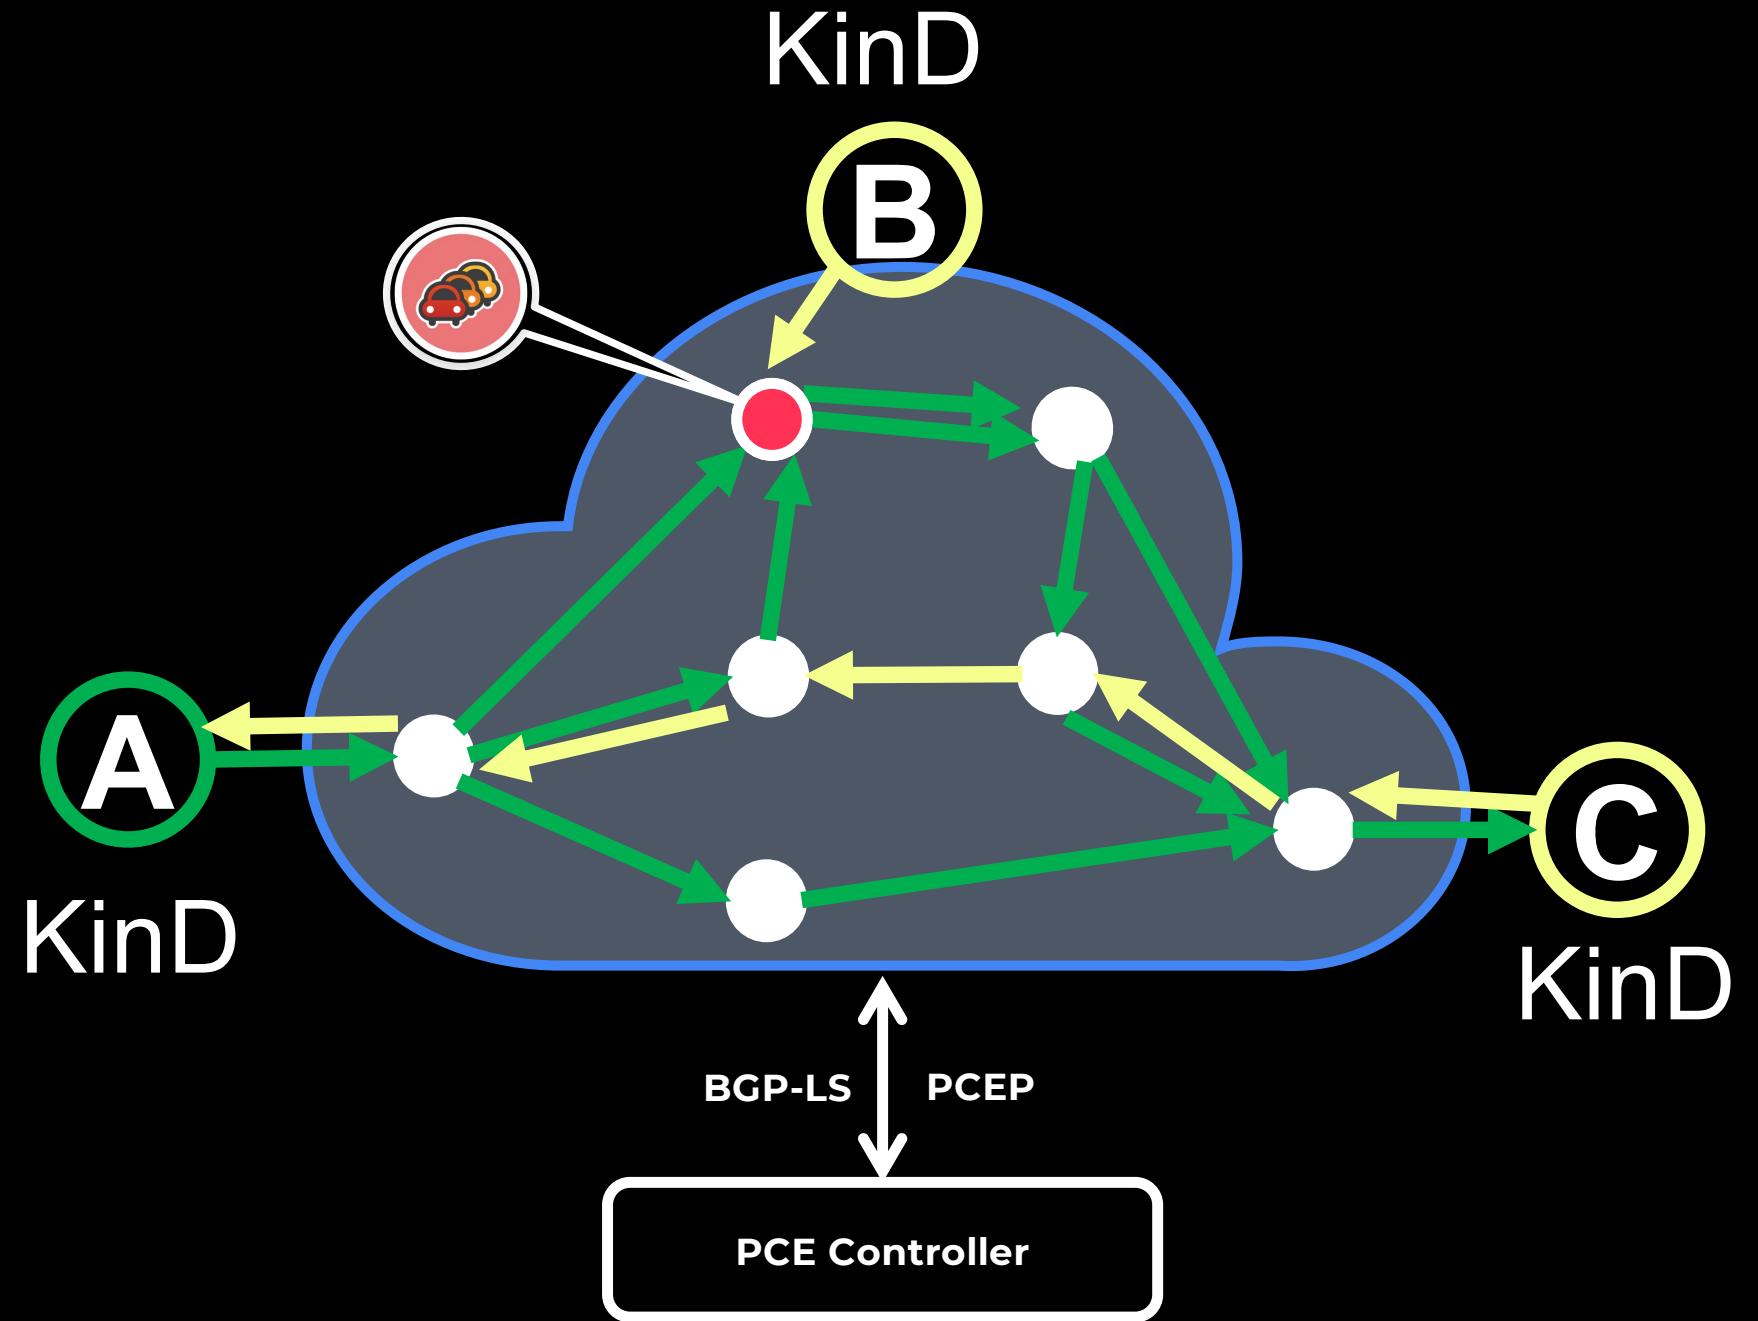
Why Use Kind + Containerlab?

- 🧪 Test **real networks** + **Kubernetes workloads** in a unified lab
- 🔗 **Integration** enables direct links between K8s pods and underlay nodes
 - Defined in a single **YAML** file for automated deployment *
- 🔄 Useful for network function testing, CNFs, and automation labs

Clab+KinD: Components & Topology

- ◆ k8s-kind nodes: create and manage Kind clusters
- ◆ ext-container nodes: expose control-plane & worker nodes for linking
 - 💡 Use exec to configure interfaces like eth1 on K8s nodes
 - 🔗 Define links between K8s nodes and network devices (e.g., SR Linux)
 - 📁 startup-config: custom Kind YAML config for multi-node clusters
 - 📡 Each Kind node joins the Docker network and underlay topology

KinD:
Simulate
Traffic
thru the
core



gNMIc

YANG

Protobuf



gRPC

HTTP2

TLS/TCP

Encodes YANG content
Encodes the gNMI operations

4 Operations:
- Capabilities
- Get
- Set
- Subscribe

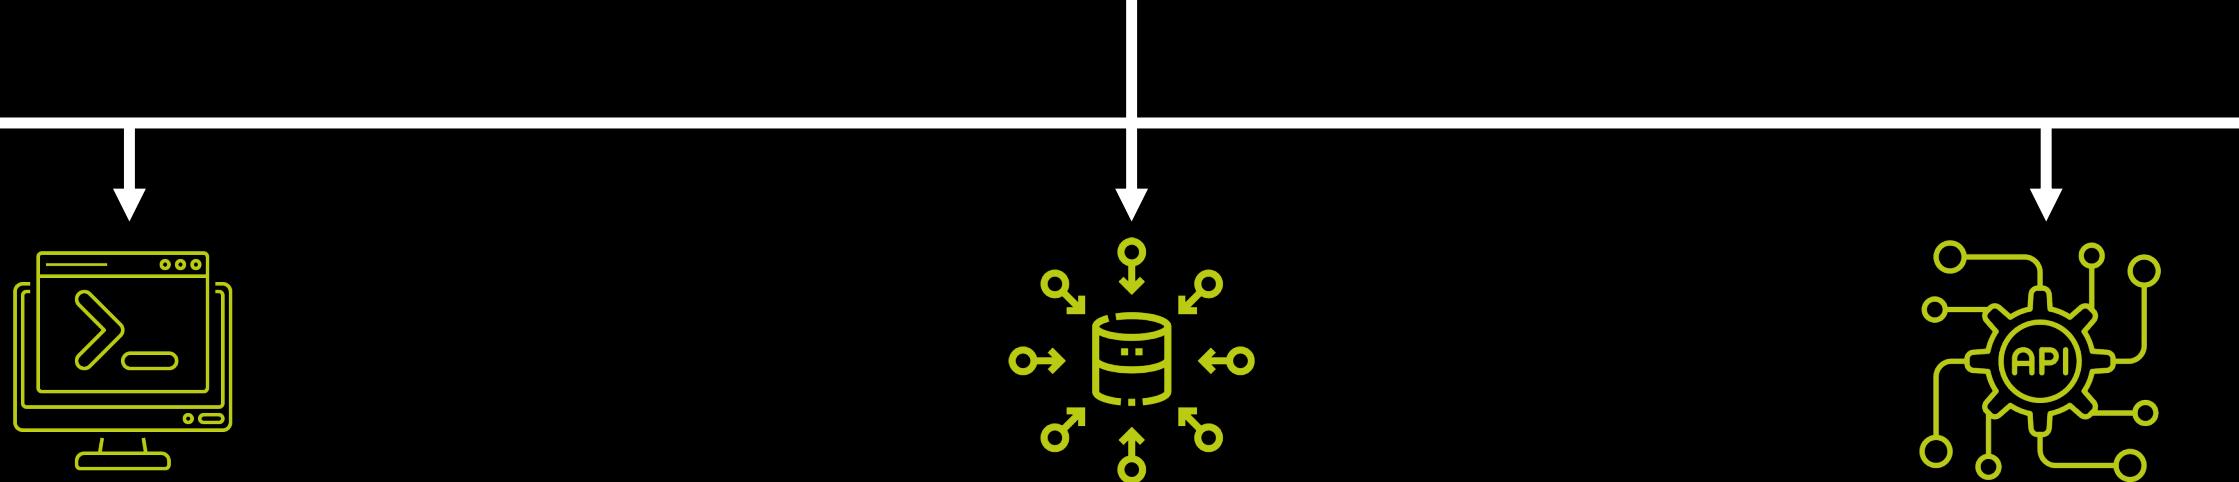
Client-Server RPC
framework

gRPC builds on HTTP2
features to multiplex
multiple streams over a
single TCP connection

gNMIc: GNMI Service Definition

- **Capabilities**: The client requests the target's capabilities (supported YANG models and their revision date)
- **Get**: The client retrieves a snapshot of the data identified by a path from the target
- **Set**: The client modifies the configuration of the target
- **Subscribe**: The client can subscribe to data identifies a set of paths, the target will stream back the data to the client periodically or when it changes.

gNMIC to gNMI

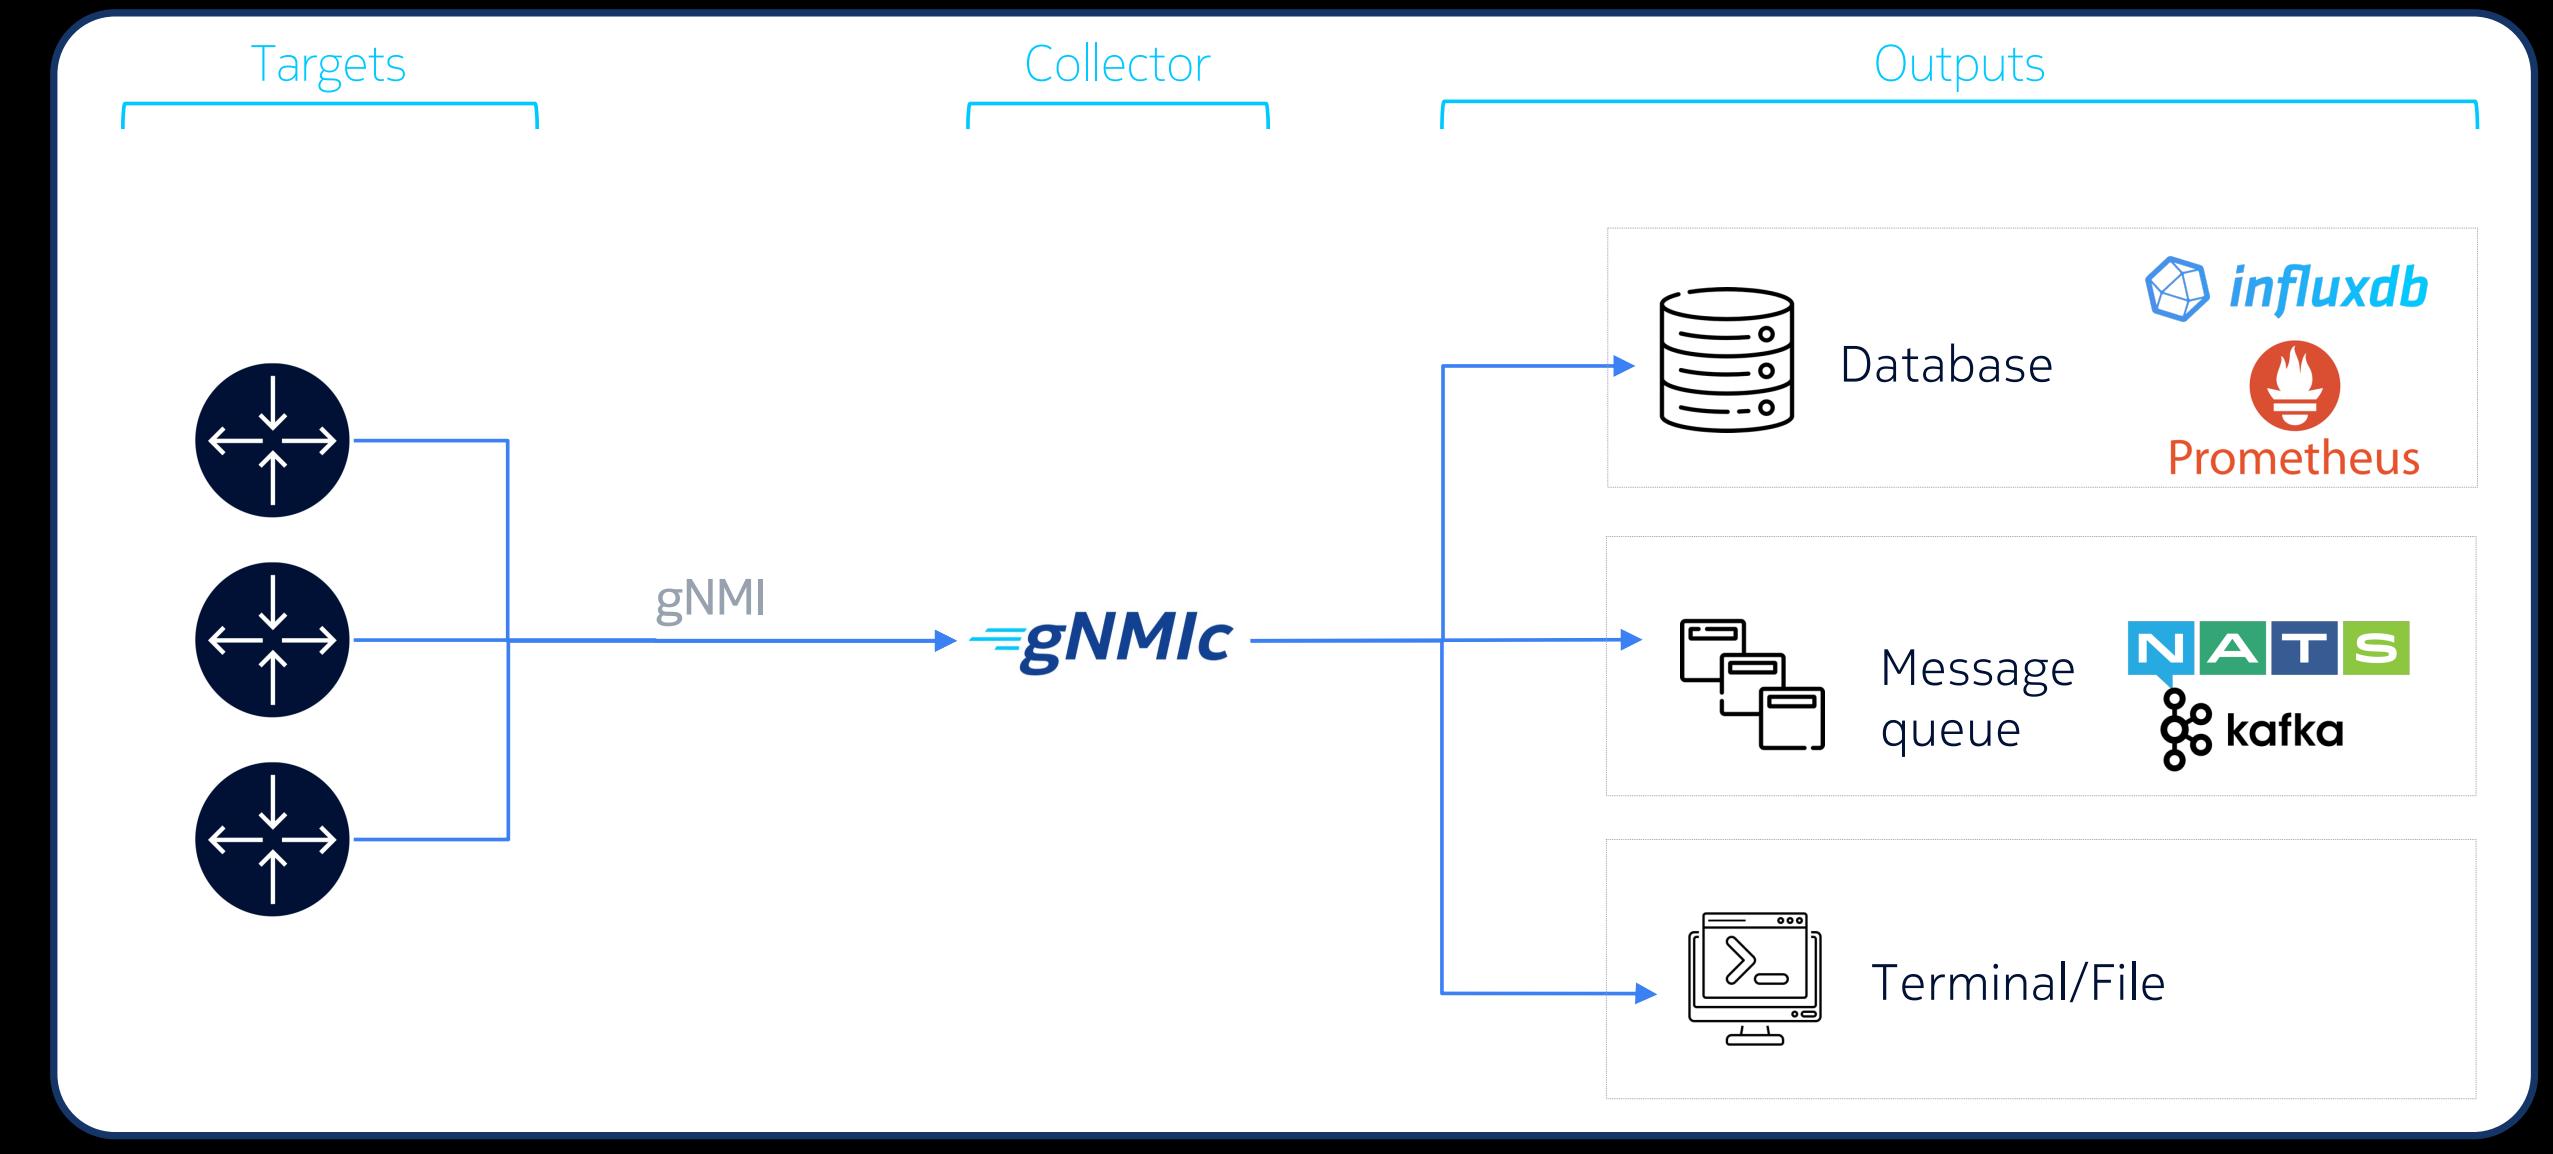


Command Line
Interface
for gNMI

Highly
Available
and flexible
collector

Go API for
gNMI with a
human touch

gNMI collector

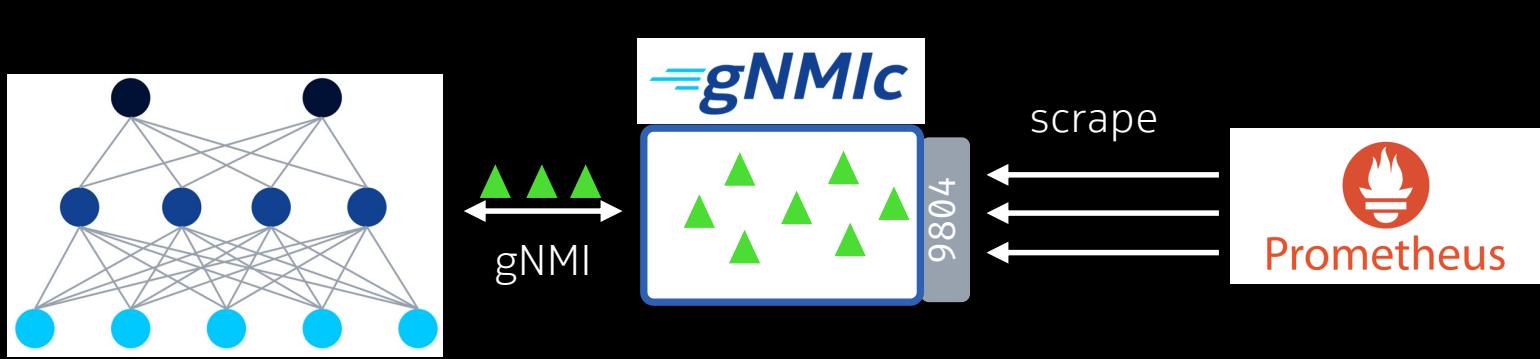


gNMI collector

More details of Prometheus on
the second part of the
Workshop!

```
outputs:  
  - output1:  
      type: prometheus  
      listen: ":9804"  
      export-timestamps: false  
      strings-as-labels: false
```

```
targets:  
  leaf1:  
    outputs:  
      - output1
```



▲ Collected telemetry metric

iperf



iperf: what?

- A network testing tool to measure bandwidth, loss, and latency.
- Focuses on TCP, UDP, and SCTP traffic performance.
- Designed for active testing between a client and a server.
- Supports modern features like JSON output and multi-threading.
- Common in performance tuning, validation, and benchmarking.

iperf: Why?

- Simple CLI interface, fast setup.
- Supports custom test durations, window sizes, and parallel streams.
- Useful for verifying network speed or identifying bottlenecks.
- Provides real-time throughput metrics.
- Ideal for both LAN and WAN environments.

iperf: Key features

- Client-server architecture (one side runs -s, the other -c).
- Customizable port, protocol, and test duration.
- Compatible with IPv4 and IPv6.
- Reports throughput, jitter, packet loss (UDP).
- Output can be human-readable or machine-parsed (JSON).

iperf: example

- s: Start as server
- c: Run as client
- p: Specify TCP/UDP port

Ensure firewall allows chosen port

Useful for parallel test sessions
or avoiding conflicts

```
# Server (listening on port 5202)
iperf3 -s -p 5202

# Client (connect to server on port 5202)
iperf3 -c 192.168.1.10 -p 5202
```

iperf in the lab

- Wrapping up iperf in a container
- Managing multiple instances to different ports
- Using k8s for orchestration
- Declarative

A black and white photograph of a large audience seated in rows of theater-style chairs, facing towards the right side of the frame where a stage would be. The lighting is dramatic, with strong highlights on the audience's faces and shadows in the auditorium.

Preparing an App for Container Deployment

Why Dockerfile Matters (vs Manual Setup)

Why Dockerfile Matters (vs Manual Setup)

- **Automates builds** – no need to set up environment manually
- **Ensures consistency** – same setup across all machines and teams
- **Version-controlled** – stored in Git, tracks changes over time
- **Portable & repeatable** – build once, run anywhere
- **Without Dockerfile** – manual setup = error-prone and hard to reproduce

Dockerfile: Key Components

- `FROM` - Base image to build your container from (e.g., `python:3.11-slim`)
- `WORKDIR` - Sets the working directory inside the container
- `COPY` - Copies files from your local machine to the container
- `RUN` - Executes shell commands during the image build process
- `CMD` / `ENTRYPOINT` - Defines default command(s) when container starts
- `EXPOSE` - Documents the port the container listens on (optional)



Why Prepare Apps for Containers?

- ✓ **Environment variables** make the app configurable (no hardcoded values)
- ✓ **Volumes** allow persistent data and external config files
- ↻ **Separation of config from code** – easier updates and reuse
- 📦 **Portability** – run the same container anywhere with different settings
- 🔒 **Secrets & credentials** managed securely via env or mounted files
- 🔨 **Dev → Test → Prod** flexibility without changing the app logic

Prep your App

- Python:

```
parser = argparse.ArgumentParser(description='Kafka Alert Exporter to Prometheus')
parser.add_argument('--bootstrap-servers', required=True, type=str, help='Kafka Bootstrap Server')
parser.add_argument('--cert', type=str, help='CA certificate path for Kafka')
parser.add_argument('--port', type=str, default="8001", help='HTTP server port (default: 8001)')
parser.add_argument('--config', required=True, type=str, help='YAML Alert topics and metrics')
parser.add_argument('--debug', action='store_true', help='Activate debug mode')
```

- Dockerfile:

```
ENV KAFKA_BOOTSTRAP_SERVERS=localhost:9092
ENV METRICS_PORT=8001
ENV ALARMS_FILE=/app/alarms.yml

CMD ["python3", "./kafka-alert-exporter.py", "--bootstrap-servers", "${KAFKA_BOOTSTRAP_SERVERS}", "--port", "${METRICS_PORT}", "--alarms", "${ALARMS_FILE}"]
```

Key Concepts for Preparation

-  Environment Variables
 - Who does not love hardcoded passwords in a Github Repo?

```
grafana:  
  image: grafana/grafana:latest  
  container_name: kafka-alert-grafana  
  volumes:  
    - ./grafana/datasources → Grafana data sources  
    - ./grafana/dashboards → Grafana dashboards  
  environment:  
    - Admin password = secret  
  depends_on:  
    - prometheus  
  ports:  
    - 3000:3000  
  restart: unless-stopped
```

Key Concepts for Preparation

-  Tagging the Image
 - Pick a donut, any donut. Surprise bugs included!



Docker Compose vs Manual Docker

- ✓ **Compose:** defines multi-container apps in one YAML file
- ✓ **Single command deploy** – docker-compose up starts everything
- ✓ **Easier networking** – containers auto-connect via service names
- ✓ **Environment config** – use .env files for flexible settings
- ✗ **Manual docker run** – repetitive, error-prone, hard to scale

Docker Compose Code Modularity

```
1  FROM ubuntu:22.04
2
3  # Install dependencies
4  RUN apt-get update && \
5      apt-get install -y wget gnupg2 curl software-properties-common
6
7  # Install Prometheus
8  RUN useradd --no-create-home --shell /bin/false prometheus && \
9      mkdir /etc/prometheus /var/lib/prometheus && \
10     cd /tmp && \
11     wget https://github.com/prometheus/prometheus/releases/download/v2.52.0/prometheus-2.52.0.linux-amd64.tar.gz && \
12     tar xvf prometheus-2.52.0.linux-amd64.tar.gz && \
13     cd prometheus-2.52.0.linux-amd64 && \
14     cp prometheus promtool /usr/local/bin/ && \
15     cp -r consoles console_libraries prometheus.yml /etc/prometheus
16
17 # Install Grafana
18 RUN wget -q -O - https://packages.grafana.com/gpg.key | apt-key add - && \
19     add-apt-repository "deb https://packages.grafana.com/oss/deb stable main" && \
20     apt-get update && \
21     apt-get install -y grafana
22
23 # Expose Ports
24 EXPOSE 9090 3000
25
26 # Start both services
27 CMD service grafana-server start && \
28     prometheus --config.file=/etc/prometheus/prometheus.yml
```

Docker Compose Code Modularity

```
1  version: '3.8'
2
3  services:
4    prometheus:
5      image: prom/prometheus:latest
6      container_name: prometheus
7      ports:
8        - "9090:9090"
9      volumes:
10        - ./prometheus/prometheus.yml:/etc/prometheus/prometheus.yml
11
12  grafana:
13    image: grafana/grafana:latest
14    container_name: grafana
15    ports:
16      - "3000:3000"
17    environment:
18      - GF_SECURITY_ADMIN_USER=admin
19      - GF_SECURITY_ADMIN_PASSWORD=admin
20    volumes:
21      - grafana-storage:/var/lib/grafana
22
23  volumes:
24    grafana-storage:
```



Docker Compose: YAML Components

- services - Main section to define each container
- build / image - Either build from Dockerfile or pull an image
- ports - Maps container ports to host ports ("8080:80")
- volumes - Mounts host paths or named volumes
- environment - Sets environment variables inside the container

Health Checks & Dependencies in Docker

- HEALTHCHECK (in Dockerfile) - tests if the container is still working
 - Example: `HEALTHCHECK CMD curl --fail http://localhost:80 || exit 1`
- Helps monitor service status for orchestrators (e.g., Compose, Swarm)
- Dependencies handled in Docker Compose with depends_on
 - Example: `web depends_on: [db, redis]`
- depends_on controls start order, but not readiness
 - Use health checks + restart policies for full dependency control



Why We Need Container Registries

- ✓ **Central storage** – push and pull images from a shared location
- ✓ **Versioning** – tag and manage image versions easily
- ✓ **Team collaboration** – share images across dev, test, and prod
- ✓ **CI/CD integration** – automate image builds and deployments
- ✓ **Security & auditing** – scan images, enforce access control
- ✗ **Without a registry** – sharing = manual, error-prone, and insecure

How to manage Container Images

- Image built and tested – now what?
-  Public Repositories
 - Push it to share and deploy anywhere
 - E.g. Docker Hub
-  Private Registries
 - Push it to your teams private registry for secure access



Push Image to Local Docker Registry

- 1. **Tag the image** with the local registry address:
 - `docker tag my-app:latest localhost:5000/my-app:latest`
 - 2. **Push it to the local registry**:
 - `docker push localhost:5000/my-app:latest`
- This assumes your **local registry** is already running on port `5000`
(e.g., started with `docker run -d -p 5000:5000 registry:2`)



Questions?